




Graph-to-Gates: Optimization of Parallel Prefix Circuits using Deep Reinforcement Learning

Presented by **Cory Brynds, Francisco Soriano, Michael Castiglia, and John Gierlach**


Motivation

- Vast amounts of arithmetic circuits are used in CPUs and GPUs to achieve unprecedented acceleration for AI, high-performance computing, and computer graphics
- Improving the design of these arithmetic circuits has a massive impact on the performance and efficiency of these processors
- Binary adders are fundamental & extensively used throughout processors (e.g. ALUs, counters)
- Example: **NVIDIA H100**
 - 144 streaming multiprocessors (SMs) per GPU
 - 128 FP32 CUDA cores and 4 tensor cores per SM
 - **18,432 FP32s CUDA cores and 576 tensor cores** per H100




The importance of design optimizations scales with core count

Review: Parallel Prefix Addition



Pre-calculate p_i, g_i signals
based on adder inputs




Compute c_i for each stage


Compute $s_i = p_i \oplus c_i$

Prefix Nodes


- Here we define an operator where the inputs are two 2-bit vectors (g_i, p_i) used to calculate the following bits
- **Generate_{out}**
 - $g_{out} = g_{in1} \vee (g_{in2} \wedge p_{in1})$
- **Propagate_{out}**
 - $p_{out} = p_{in1} \wedge p_{in2}$
- The output is a 2-bit vector (g_{out}, p_{out}) that represents the combined generate and propagate signals from both input signals



Designing a Parallel Prefix Adder




Sequential Carry (RCA)




Parallel Carry

8-Bit Parallel Prefix Adder



Serial (RCA)



Parallel (Sklansky)

Adder Type	Delay (Stages)	Area (Node count)
Serial (RCA)	7	7 (7 OR, 14 AND gates)
Parallel (Sklansky)	$\log_2(8) = 3$	12 (12 OR, 24 AND gates)

The Parallel Prefix Problem

- Existing parallel prefix algorithms tend to offer minimum delay **OR** reduced area but do not offer a balanced tradeoff between the two metrics
- There are $\frac{(N-1)(N-2)}{2}$ locations where a prefix node may or may not exist, meaning there are $O(2^{N^2})$ valid N-input prefix graphs!
- Exhaustively searching for the best tradeoff in this enormous design space is impractical
 - The total design space for valid prefix graphs for a 32-bit adder is $O(2^{465})$

Is it possible for a model to explore and optimize graphs within this enormous design space?

Existing Prefix Graph Optimization Solutions

- **Historical Parallel Prefix Algorithms**

- Serial (RCA): **Worst delay**, best area, minimal fanout
- Sklansky (1960): Best delay, higher area than serial, **worst fanout**
- Kogge-Stone (1973): Best delay, **worst area**, minimal fanout
- Brent-Kung (1982): Higher delay than Sklanksy/KS, lower area than Sklansky, bounded but not minimal fanout

Algorithm	Delay (# of stages)	Area (# of nodes)	Max fanout
Serial	$n - 1$	$n - 1$	2
Sklansky	$\log_2 n$	$\frac{1}{2} n \log_2 n$	$\frac{1}{2} n$
Brent-Kung	$2\log_2 n - 2$	$2n - \log_2 n - 2$	$\log_2 n$
Kogge-Stone	$\log_2 n$	$n \log_2 n - n + 1$	2

Existing Prefix Graph Optimization Solutions

- Historical Parallel Prefix Algorithms
 - Serial (RCA): **Worst delay**, best area, minimal fanout
 - Sklansky (1960): Best delay, higher area than serial, **worst fanout**
 - Kogge-Stone (1973): Best delay, **worst area**, minimal fanout
 - Brent-Kung (1982): Higher delay than Sklanksy/KS, lower area than Sklansky, bounded but not minimal fanout
- **Analytic Approaches**
 - Pruning via heuristics and delay/area prediction with ML model
 - Simulated annealing: randomly modify prefix graph towards optimal point – using an analytic model

Existing Prefix Graph Optimization Solutions


- Historical Parallel Prefix Algorithms
 - Serial (RCA): **Worst delay**, best area, minimal fanout
 - Sklansky (1960): Best delay, higher area than serial, **worst fanout**
 - Kogge-Stone (1973): Best delay, **worst area**, minimal fanout
 - Brent-Kung (1982): Higher delay than Sklanksy/KS, lower area than Sklansky, bounded but not minimal fanout
- Analytic Approaches
 - Pruning via heuristics and delay/area prediction with ML model
 - Simulated annealing: randomly modify prefix graph towards optimal point – using an analytic model
- **PrefixRL**: RL-based approach that explores design space with physical synthesis evaluation metrics

PrefixRL: An RL Model for Optimizing Parallel Prefix Circuits

- Reinforcement learning framework that trains an agent to explore an unrestricted prefix graph design space to **co-optimize** for delay and area
 - Incorporates physical synthesis **inside of the training loop** to extract more accurate metrics than prior analytical models
 - Optimizes prefix circuits for 32-bit and 64-bit adders that **Pareto-dominate** all prior designs – offering better area-delay tradeoffs with at least 1 improved metric and no metrics worse than previous designs


Reinforcement Learning Environment

- **Task** – maximize cumulative reward via finding prefix adder graphs that optimize area & delay tradeoff
- **State space** – all valid NxN prefix graphs
- **Action** – add or delete a prefix node
- **Reward** – returns area-delay tradeoff



Deep Q-Network (DQN) Architecture

- Our Q-network utilizes a CNN architecture
- The **input** to the DQN is an $N \times N$ grid with 4 channels indicating:
 1. Presence of a node (1 if present, else 0)
 2. Deletability of a node (1 if delectable, else 0)
 3. Level of node (Normalized to [0,1])
 4. Fanout of node (Normalized to [0,1])
- The DQN **outputs** an $N \times N$ grid with 4 channels, corresponding to the Q values for adding and deleting a node, in terms of area and delay
- The RL agent iteratively constructs the prefix graph by **selecting actions** according to its policy informed by **Q-values**
- We use **B=32, C=256** for 32b and 64b adder training



Deep Q-Learning Algorithm

- To train our reinforcement learning model, we implement the classic Deep Q-Learning Algorithm¹
 1. Take an action and observe (s_i, a_i, s_{i+1}, r_i) , adding it to the replay buffer
 2. Sample a random batch (s_j, a_j, s_{j+1}, r_j) from the replay buffer
 3. Compute the bellman target $y_j = r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)$ (the sum of the current reward and the expected future reward from the **target network** $Q_{\phi'}$)
 4. Update the parameters of the **online network** based on $\text{Loss}(Q_{\phi}(s_j, a_j), y_j)$
 5. Every N steps, update $\phi' \leftarrow \phi$ (synchronize target network's parameters with online network)




Network Initialization

- In the initialization phase, the online and target Q-networks are initialized with random weights
- Separate target and online networks are used to create a separation between network evaluation and network updating to avoid feedback loops and improve error reduction while training.
- The Adam optimizer is initialized with a dynamic learning rate - exponentially decreasing as training goes on
- We use a **replay buffer** with 400000 elements to cache the results of past experiences to improve training stability. During the network update phase, batches will be randomly sampled from the replay buffer

Episodic Reset

- At the beginning of a new episode, the environment is reset to an initial graph state
- The PrefixRL paper resets to a ripple-carry adder (RCA) graph at the beginning of each episode
- Modifying the initial starting state may allow the Q-network to converge to an optimal policy more quickly¹
- For our implementation, we cyclically reset to a ripple carry, Sklansky, and Brent-Kung adder on each episode




Epsilon Greedy

- The epsilon greedy approach is used to determine if the network will randomly add or delete a node from the current prefix graph
- This is done by take current progress of the episodes that are currently completed and applying an exponential increase to the epsilon value.
- Over the course of the training, the rate at with the nodes are added or deleted will decrease.


$$\epsilon(e) = 0.4^{\left(1+7\frac{i(e)}{1023}\right)}$$

Greedy Epsilon Curve




Prefix Graph Action & Reward Function Calculation

- Once an action is selected, a node will be added or deleted from the current graph state
- To evaluate the reward within our environment, the prefix graph will be converted to a Verilog netlist
- From there, the netlist is synthesized to generate the delay and area metrics of the new parallel prefix adder
- Finally, the reward is calculated based on the delay and area reported from the synthesis tool





Evaluation Flow

- Our evaluation flow uses a suite of open-source EDA tools
- **Graph-to-netlist conversion** – generates RTL netlist based on graph produced by the RL agent
- **Yosys** – logic synthesis to produce get technology-mapped netlist (Using open-source Nangate45 standard cell library)
- **OpenROAD** – open-source place-and-route flow to evaluate circuit delay & area
(Adapted from ArithTreeRL repository)
 - Two separate flows: full flow (full placement and routing optimizations, ~10 seconds for 32b adder) and fast flow (only placement, no routing, ~1 second for 32b adder)




Analytic Model of Area and Delay

- Performing circuit synthesis each environment step can be quite time-consuming, especially without parallelization
- Training can be sped up significantly by using analytic metrics for area and delay, at the cost of accuracy
- $\text{area} = \text{prefix node count}$
- $\text{delay} = \text{delay}_{\text{internal}} + \text{delay}_{\text{driving}} + \text{delay}_{\text{sum}}$
 - $\text{delay}_{\text{internal}} = \text{stage count} + (\text{D}_I * \text{delay}_{\text{node}})$
 - $\text{delay}_{\text{driving}} = \text{D}_O * (\text{critical path fanout})$
 - $\text{delay}_{\text{sum}} = 1$ (delay of the final $p_i \oplus c_i$ stage)
 - $\text{D}_O = 0.5, \text{D}_I = 1.0^1$




Scalarized Multi-Objective Reward

- The Q-network learns the Q-values for our two objectives, area and delay, separately
- To select the optimal action (add or delete a node), these separate Q values must be scalarized according to a weight scalar w
- $r_t = w(\text{delay}_t - \text{delay}_{t+1}) + (1 - w)(\text{area}_t - \text{area}_{t+1})$
- w is a value in $[0,1]$ that allows for different tradeoffs in area and delay
 - $w = 1$ optimizes only for delay and $w = 0$ optimizes only for area
- The Q-network's goal is to find a policy to maximize r_t for a given w



Q-Network Update

- To update the parameters of the online network based on gathered experience, a batch of (s_j, a_j, s_{j+1}, r_j) is sampled from the replay buffer
- The target network is evaluated to compute the bellman target y_j for the sampled experience
- The online network's parameters are updated according to
$$\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} (Q_\phi(s_j, a_j) - y_j)$$
- Every 60 steps of the training loop of a given episode, the target network and online network are synchronized to ensure the target network has the updated weights from the online network




Decaying Learning Rate

- The learning rate of the Q-network determines how much the parameters of a model are updated each environment step
- In general, the parameters should be updated in smaller increments as the Q-network converges to an optimal policy
- A decaying learning rate will slow down the rate at which the Q-network's parameters are updated with each episode to avoid any drastic design changes near the end of training
- This is done by using a scheduler to decrease the learning rate over each episode and step


$$\text{lr}(t) = \text{lr}_0 \gamma^t$$

Decaying Learning Rate for 500 Steps



Workload Parallelization

- To ensure timely results, we parallelized the RL computations using AMD & NVIDIA GPUs on personal machines and clusters (e.g. ARCC, VLSI)
- Additionally, there were parallel processes that called Yosys and OpenROAD



Note on GPU/CPU Utilization

```
Every 0.1s: rocm-smi                                         draco-eng-A01: Sat Nov 15 13:53:09 2025


=====
ROCM System Management Interface
Concise Info
=====
Device Node IDs      Temp   Power  Partitions          SCLK    MCLK    Fan   Perf  PwrCap  VRAM%  GPU%
^[3m          (DID,     GUID) (Edge) (Avg)  (Mem, Compute, ID) ^[0m
=====
0    1    0x744c, 4672  54.0°C 51.0W  N/A, N/A, 0    42Mhz 772Mhz 14.9% auto 339.0W 71%  0%
1    2    0x164e, 18139 43.0°C 22.131W N/A, N/A, 0    N/A   2400Mhz 0%   auto N/A   3%   0%
=====
End of ROCm SMI Log
=====

0[     0.0%] 4[     0.0%] 8[     0.0%] 12[    0.0%] 16[    0.0%] 20[    0.0%] 24[    0.0%] 28[    0.0%]
1[||||| 66.7%] 5[||| 37.5%] 9[     0.0%] 13[    0.0%] 17[||| 20.0%] 21[    0.0%] 25[    0.0%] 29[    0.0%]
2[||| 12.5%] 6[     0.0%] 10[    0.0%] 14[    0.0%] 18[    0.0%] 22[    0.0%] 26[    0.0%] 30[    0.0%]
3[     0.0%] 7[     6.2%] 11[    0.0%] 15[    0.0%] 19[    0.0%] 23[    6.7%] 27[    0.0%] 31[    0.0%]
Mem[||||| 33.1G/93.4G] Tasks: 352, 3297 thr; 2 running
Swp[ ]           1.25M/2.00G Load average: 11.36 8.11 5.36
Uptime: 24 days, 02:39:59

PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
2810535 root  20  0 31.0G 4289M 2329M R 83.0 4.4 6:43.52 python3 graph_to_gates.py -n 32 --adder type 0 -b 32 --w scalar 0.7 --output dir adder_32b_batch16_weight0p7.st
2863089 mcastigli 20  0 26456 9984 3840 R 38.3 0.0 0:32.43 htop -d 1
485070 mcastigli 20  0 28.2G 761M 135M S 6.4 0.8 2:37.56 /snap/firefox/7084/usr/lib/firefox/firefox -contentproc -isForBrowser -prefsHandle 0:44657 -prefMapHandle 1:276
622705 mcastigli 20  0 2755M 317M 108M S 6.4 0.3 0:21.64 /snap/firefox/7084/usr/lib/firefox/firefox -contentproc -isForBrowser -prefsHandle 0:44658 -prefMapHandle 1:276
1329803 mcastigli 20  0 8458M 474M 185M S 6.4 0.5 21:10.91 /usr/bin/gnome-shell
2740348 mcastigli 20  0 28.2G 761M 135M S 6.4 0.8 0:01.36 /snap/firefox/7084/usr/lib/firefox/firefox -contentproc -isForBrowser -prefsHandle 0:44657 -prefMapHandle 1:276
1 root    20  0 164M 12544 7680 S 0.0 0.0 2:11.07 /sbin/init splash
```

Parallelization Optimization Results

Training Time for Various Levels of Parallelization




Graph to Gates User Interface

- Input arguments printed to screen:
 - Input adder bit-width/type
 - Use of analytic model or physical-aware model
 - Weight scalar (to incentivize area or delay)
 - Training hyperparameters




ARGUMENT CONFIGURATION	
Input bitwidth:	32
Adder type:	RCA
Use analytic model:	True
Number of training steps per episode:	5000
Number of training episodes:	100
Weight scalar for area and delay:	0.5
Batch size:	192
OpenROAD path:	../OpenROAD/prefix-flow/
Flow type:	fast_flow
Output directory:	out/
Save Verilog:	False
Parallel evaluation:	Enabled
Restore from:	None
Disable checkpointing:	False


32b Adder Training Results for Different Weight Scalars



$w = 0.3$




$w = 0.5$




$w = 0.7$

32b Adder Results (Trained with w=0.3)




64b Adder Training Results




$$w = 0.7$$

64b Adder Results (Trained with w=0.7)



32b Adder Results: Analytic Model

- Since the analytic model for delay computes only the critical path, it does not have context for the fanout of most intermediate nodes
- This leads to non-ideal graph modifications



Evaluation Against Commercial Tools



- Our RL-optimized adders were evaluated using open-source EDA tools and standard cell libraries
- To see how our results generalize to commercial EDA tools and cell libraries, we evaluate results with Synopsys RTL Architect using a 32nm process technology
- In addition, we measure the power consumption using PrimePower
- For the following results, we compare our RL-generated adders against Synopsys DesignWare adders

```
module adder_top #(parameter N = 32) (
    input [N-1:0] a,
    input [N-1:0] b,
    output [N-1:0] s,
    output cout
);



    assign {cout, s} = a + b;

endmodule
```

32b Evaluation Against Commercial Tools



64b Evaluation Against Commercial Tools



Project Limitations

Compute, Compute, Compute

- Using the exact parameters from PrefixRL for 32b adder training (5000 steps/episode, 100 episodes, batch size of 192), full training would take 150 CPU days sequentially
- To obtain our results, we used a much-reduced parameter set (500 steps/episode, 10 episodes, batch size of 16/32)

Power Optimization

- The training process does **not** optimize for power – minimizing delay and/or area *may* increase power
- Without multi-objective optimization including power, the RL agent may produce designs that have balanced delay and area but not power-efficiency

Signal Arrival Times

- We assume that all input signals arrive to the graph with a uniform latency
- In reality, input bits may arrive with non-uniform latencies, requiring different optimizations



Thank You

Questions?